

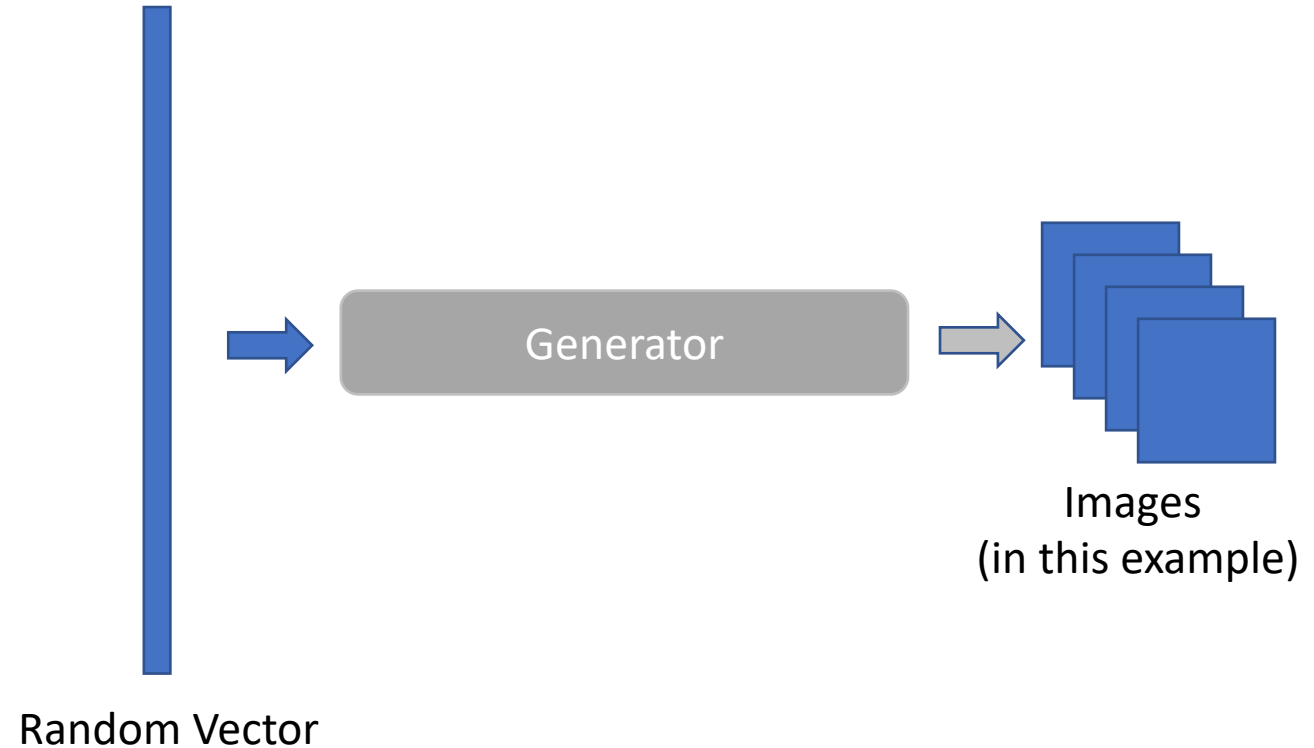
# Discriminator

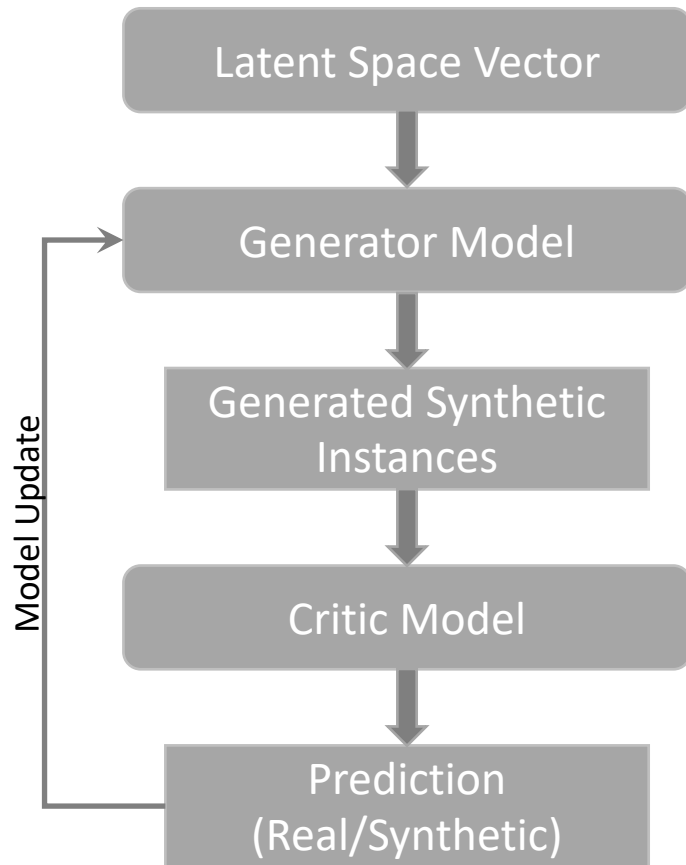
- Objective: Learn to discriminate between real & fake images.
- Supervised training, Classification model.
- Input: Images  
Output: Labels (real/fake)  
Loss: Binary Cross Entropy



# Generator

- Objective: Generate fake images...that should **“fool”** the discriminator.  
Ergo, Adversarial.
- Input: Random Vector (Diversity/variability)  
Output: Images



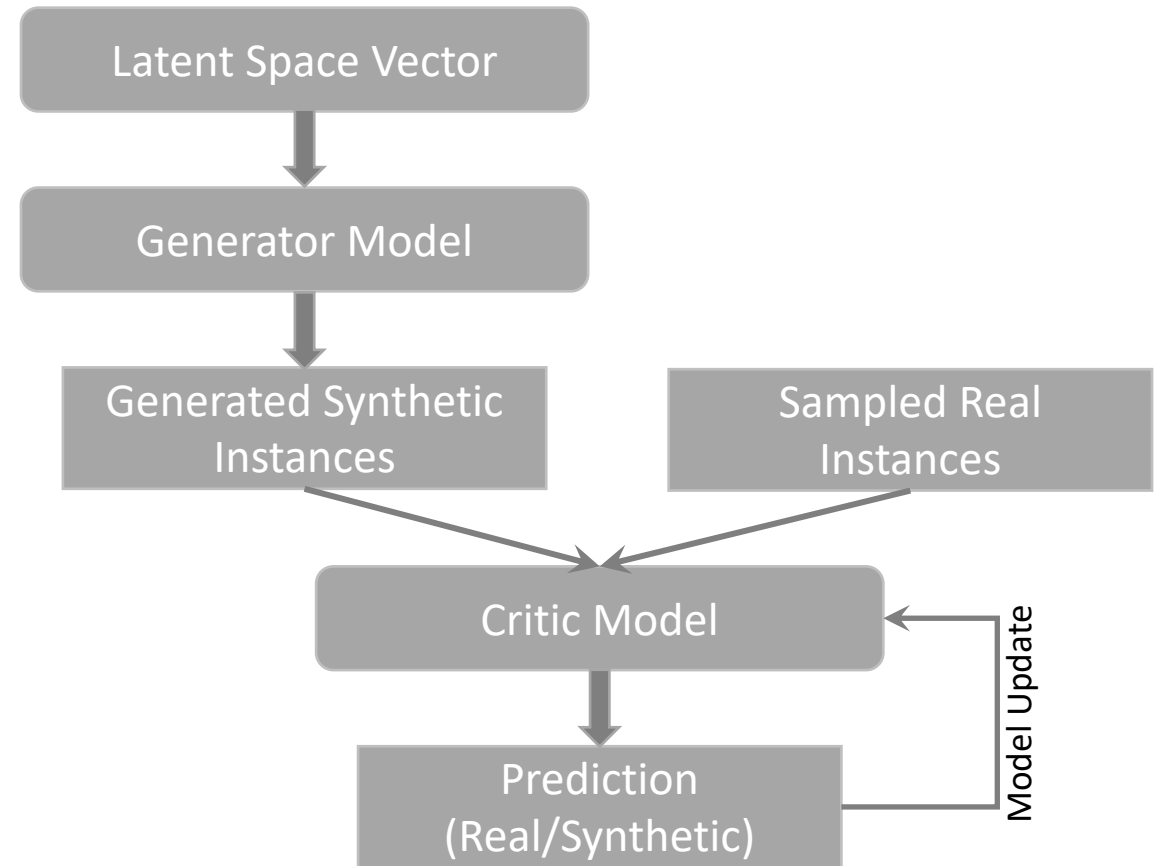


# Generator Training

- Ideal answer for the generator is all images were real.
- Hence target vector =  $[1,1,1,1,1,\dots,1]$
- Discriminator Prediction:  $[0,1,0,1,1,0,0,\dots,0]$
- Calculate the BCE loss between target and prediction.
- Apply gradient descent...but *only on the parameters of the generator model*.

# Discriminator Training

- Old fashioned supervised training step.



GAN Training = Discriminator Train\* + Generator Train\*

